

Assignment -3

Python Programming

Assignment Date	07 October 2022
Student Name	Ms.Pavithra.C
Student Roll Number	11329071028
Maximum Marks	2 Marks

Problem Statement: Abalone Age Prediction

import modules

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import scale
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import
accuracy_score,confusion_matrix,classification_report
```

convert csv file to dataframe

```
ns=pd.read_csv(r"C:\Users\IBM\Documents\abalone.csv")
```

```
ns.head()
```

```
Sex Length Diameter Height Whole weight Shucked weight Viscera weight
\
0   M   0.455    0.365   0.095      0.5140       0.2245     0.1010
1   M   0.350    0.265   0.090      0.2255       0.0995     0.0485
2   F   0.530    0.420   0.135      0.6770       0.2565     0.1415
3   M   0.440    0.365   0.125      0.5160       0.2155     0.1140
4   I   0.330    0.255   0.080      0.2050       0.0895     0.0395
```

```
Shell weight Rings
0        0.150     15
1        0.070      7
2        0.210      9
3        0.155     10
4        0.055      7
```

```

ns.tail()

      Sex  Length  Diameter  Height  Whole weight  Shucked weight \
4172    F    0.565     0.450    0.165      0.8870      0.3700
4173    M    0.590     0.440    0.135      0.9660      0.4390
4174    M    0.600     0.475    0.205      1.1760      0.5255
4175    F    0.625     0.485    0.150      1.0945      0.5310
4176    M    0.710     0.555    0.195      1.9485      0.9455

      Viscera weight  Shell weight  Rings
4172          0.2390      0.2490     11
4173          0.2145      0.2605     10
4174          0.2875      0.3080      9
4175          0.2610      0.2960     10
4176          0.3765      0.4950     12

```

To know the shape of dataset

```
ns.shape
```

```
(4177, 9)
```

To calculate the age add '1.5' with Rings

```

age = []
for h in ns["Rings"]:
    age.append(h+1.5)

ns['Age'] = age

ns['Age']

0      16.5
1      8.5
2     10.5
3     11.5
4      8.5
...
4172   12.5
4173   11.5
4174   10.5
4175   11.5
4176   13.5
Name: Age, Length: 4177, dtype: float64

ns.drop(columns=['Rings'], axis=1, inplace=True)

ns.head()

      Sex  Length  Diameter  Height  Whole weight  Shucked weight  Viscera weight \

```

```

0   M   0.455      0.365  0.095      0.5140      0.2245  0.1010
1   M   0.350      0.265  0.090      0.2255      0.0995  0.0485
2   F   0.530      0.420  0.135      0.6770      0.2565  0.1415
3   M   0.440      0.365  0.125      0.5160      0.2155  0.1140
4   I   0.330      0.255  0.080      0.2050      0.0895  0.0395

  Shell weight  Age
0       0.150  16.5
1       0.070  8.5
2       0.210  10.5
3       0.155  11.5
4       0.055  8.5

ns.tail()

      Sex  Length  Diameter  Height  Whole weight  Shucked weight \
4172   F    0.565     0.450   0.165      0.8870      0.3700
4173   M    0.590     0.440   0.135      0.9660      0.4390
4174   M    0.600     0.475   0.205      1.1760      0.5255
4175   F    0.625     0.485   0.150      1.0945      0.5310
4176   M    0.710     0.555   0.195      1.9485      0.9455

  Viscera weight  Shell weight  Age
4172        0.2390      0.2490  12.5
4173        0.2145      0.2605  11.5
4174        0.2875      0.3080  10.5
4175        0.2610      0.2960  11.5
4176        0.3765      0.4950  13.5

```

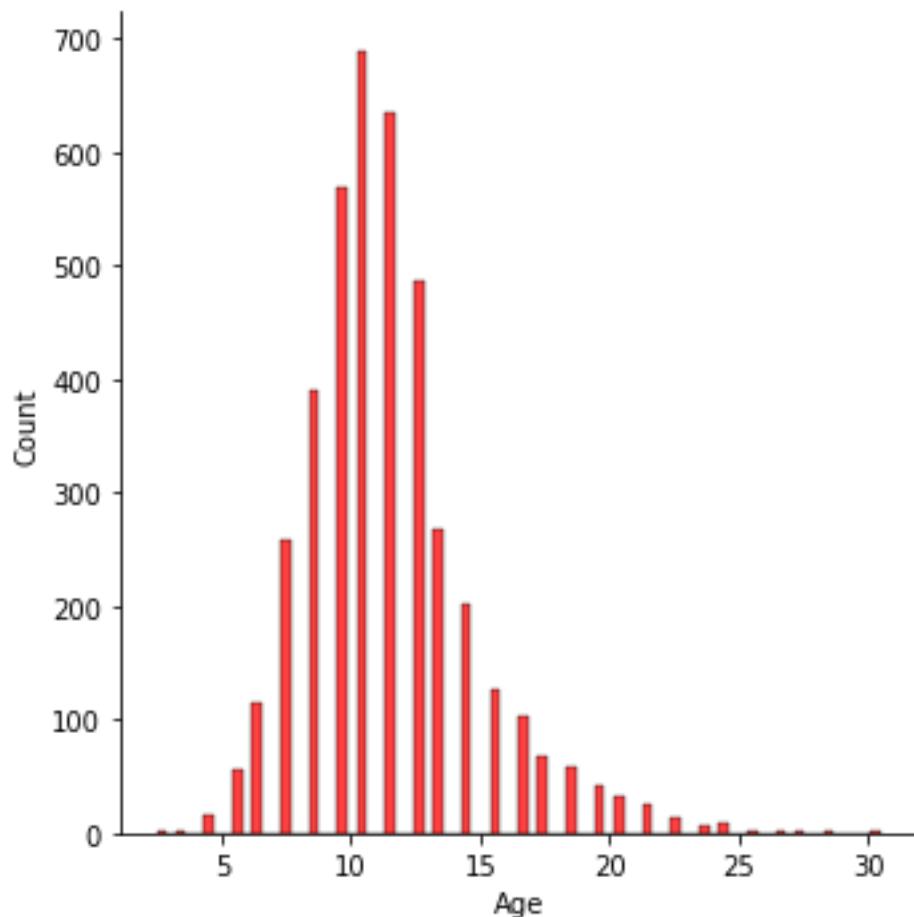
Perform Visualizations

Univariate Analysis

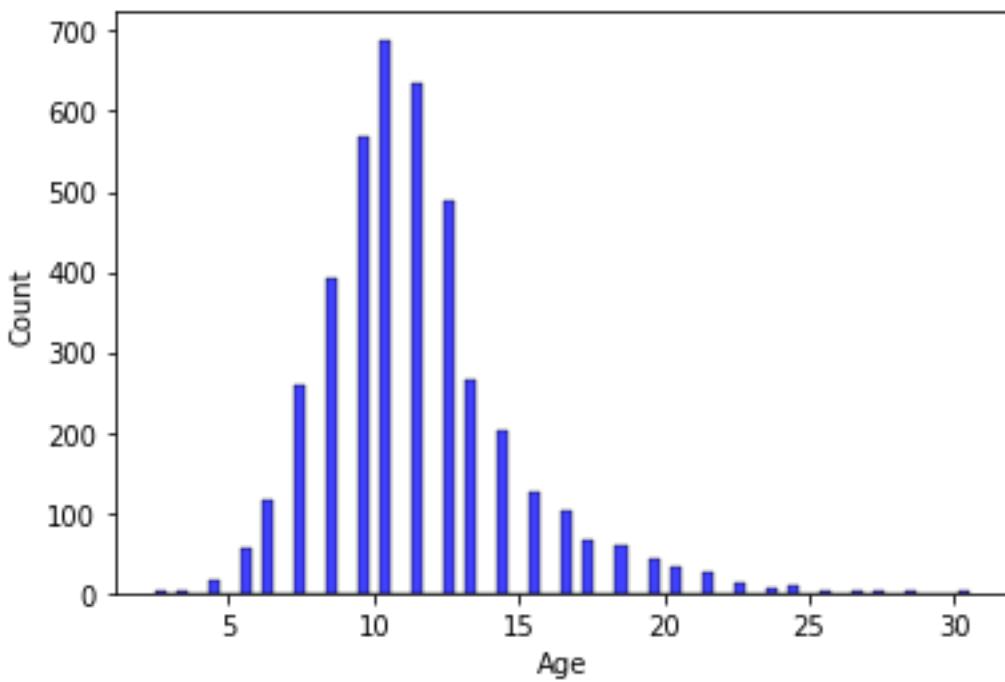
```

xz.displot(ns["Age"], color='red')
<seaborn.axisgrid.FacetGrid at 0x131100d0c18>

```




```
xz.histplot(x=ns['Age'],color='blue')  
<AxesSubplot:xlabel='Age', ylabel='Count'>
```



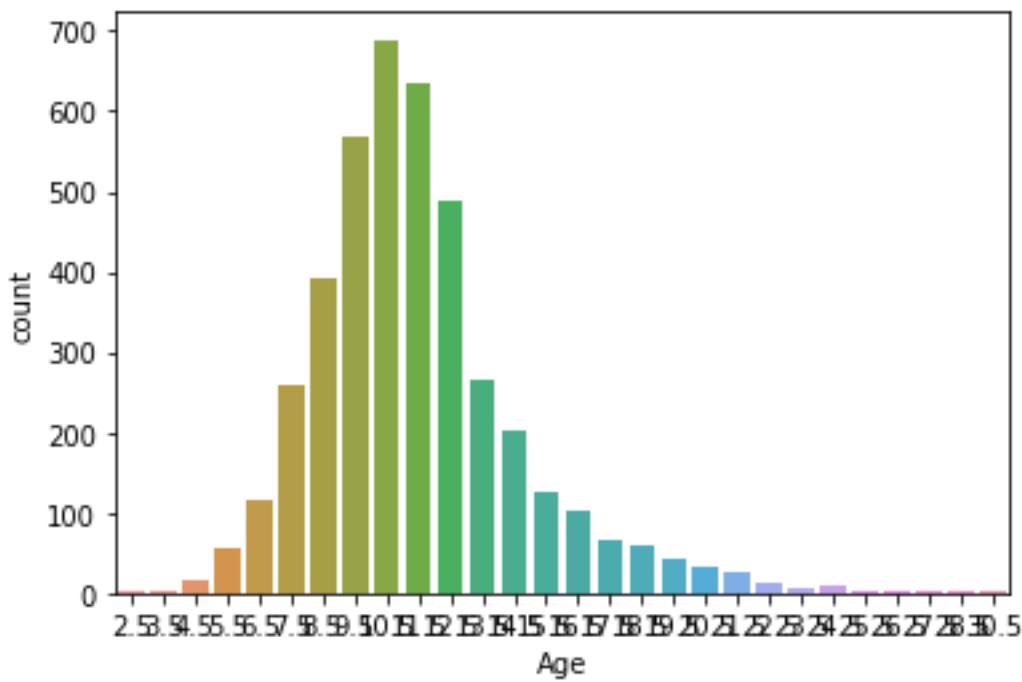
```
xz.boxplot(x=ns['Age'],color='red')
```

```
<AxesSubplot:xlabel='Age'>
```




```
xz.countplot(x=ns['Age'])
```

```
<AxesSubplot:xlabel='Age', ylabel='count'>
```




Bi-Variate Analysis

```
xz.barplot(x=ns['Height'],y=ns['Age'])  
<AxesSubplot:xlabel='Height', ylabel='Age'>
```




```
xz.lineplot(x=ns['Age'],y=ns['Height'], color='red')
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```



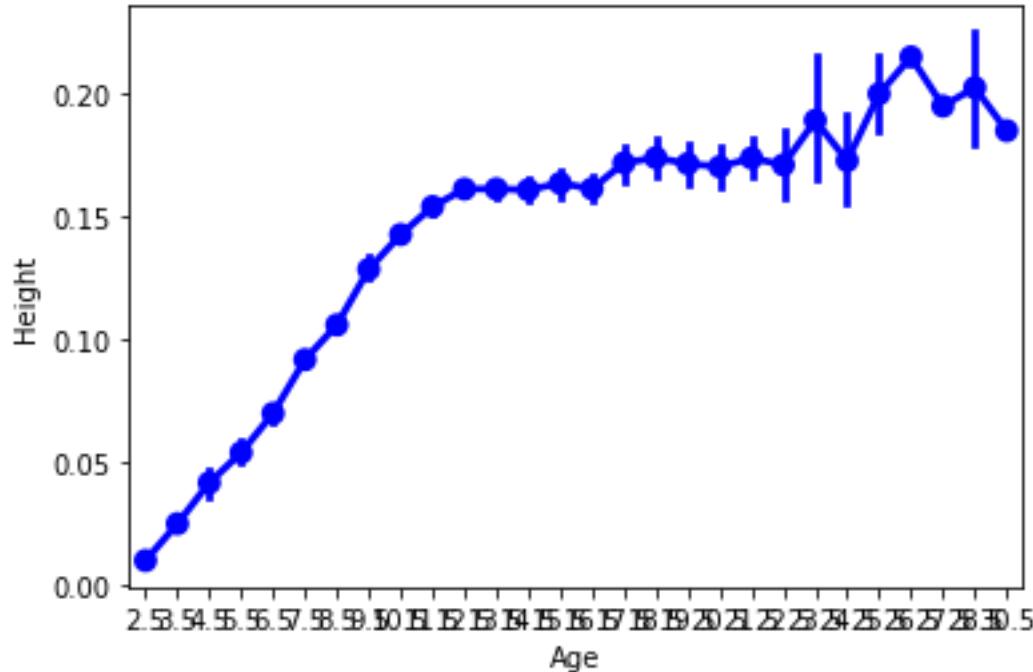
```
xz.scatterplot(x=ns['Age'],y=ns['Height'],color='blue')
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```




```
xz.pointplot(x=ns['Age'], y=ns['Height'], color="blue")
```

```
<AxesSubplot:xlabel='Age', ylabel='Height'>
```




```
xz.regplot(x=ns['Age'],y=ns['Height'],color='Red')  
<AxesSubplot:xlabel='Age', ylabel='Height'>
```



Multi-Variate Analysis

```
xz.pairplot(data=ns[["Height","Length","Diameter","Age","Whole  
weight","Shucked weight","Viscera weight","Shell weight"]])
```

<seaborn.axisgrid.PairGrid at 0x13110619438>



Perform descriptive statistics on the dataset

ns.describe()

	Length	Diameter	Height	Whole weight	Shucked weight	\
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	
std	0.120093	0.099240	0.041827	0.490389	0.221963	
min	0.075000	0.055000	0.000000	0.002000	0.001000	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	
max	0.815000	0.650000	1.130000	2.825500	1.488000	

	Viscera weight	Shell weight	Age
count	4177.000000	4177.000000	4177.000000
mean	0.180594	0.238831	11.433684
std	0.109614	0.139203	3.224169
min	0.000500	0.001500	2.500000
25%	0.093500	0.130000	9.500000
50%	0.171000	0.234000	10.500000
75%	0.253000	0.329000	12.500000
max	0.760000	1.005000	30.500000

check the null value and if exist delete the null value

```
data.isnull().sum()
```

```
Sex          0
Length       0
Diameter     0
Height        0
Whole weight  0
Shucked weight 0
Viscera weight 0
Shell weight   0
Rings         0
dtype: int64
```

Find the outliers and replace them outliers

```
outliers=ns.quantile(q=(0.25,0.75))
```

outliers

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight
\						
0.25	0.450	0.35	0.115	0.4415	0.186	0.0935
0.75	0.615	0.48	0.165	1.1530	0.502	0.2530

	Shell weight	Age
0.25	0.130	9.5
0.75	0.329	12.5


```
d1 = ns.Age.quantile(0.25)
d2 = ns.Age.quantile(0.75)
r = d2 - d1
lower_limit = d1 - 1.5 * r
ns.median(numeric_only=True)
```

Length	0.5450
Diameter	0.4250
Height	0.1400
Whole weight	0.7995

```
Shucked weight      0.3360
Viscera weight     0.1710
Shell weight       0.2340
Age                 10.5000
dtype: float64
```

```
ns['Age'] = np.where(ns['Age'] < lower_limit, 7, ns['Age'])
xz.boxplot(x=ns.Age, showfliers = False)
```

```
<AxesSubplot:xlabel='Age'>
```



Check for Categorical columns and perform encoding

```
ip = LabelEncoder()
```

```
ns.Sex = ip.fit_transform(ns.Sex)
```

```
ns.head()
```

```
   Sex  Length  Diameter  Height  Whole weight  Shucked weight \
0    2    0.455     0.365   0.095      0.5140        0.2245
1    2    0.350     0.265   0.090      0.2255        0.0995
2    0    0.530     0.420   0.135      0.6770        0.2565
3    2    0.440     0.365   0.125      0.5160        0.2155
4    1    0.330     0.255   0.080      0.2050        0.0895

   Viscera weight  Shell weight  Age
0         0.1010     0.150  16.5
1         0.0485     0.070   8.5
2         0.1415     0.210  10.5
```

```

3      0.1140      0.155  11.5
4      0.0395      0.055   8.5

```

Split the data into dependent and independent variables

```
d= ns[ "Sex" ]
```

```
d
```

```

0      2
1      2
2      0
3      2
4      1
..
4172    0
4173    2
4174    2
4175    0
4176    2

```

```
Name: Sex, Length: 4177, dtype: int32
```

```
id=ns.drop(columns=[ "Sex" ],axis=1)
id.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	\
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	

	Shell weight	Age
0	0.150	16.5
1	0.070	8.5
2	0.210	10.5
3	0.155	11.5
4	0.055	8.5

Scale the independent variables

```
id_Scaled = pd.DataFrame(scale(id), columns=id.columns)
id_Scaled.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	\
0	-0.574558	-0.432149	-1.064424	-0.641898	-0.607685	-0.726212	
1	-1.448986	-1.439929	-1.183978	-1.230277	-1.170910	-1.205221	
2	0.050033	0.122130	-0.107991	-0.309469	-0.463500	-0.356690	
3	-0.699476	-0.432149	-0.347099	-0.637819	-0.648238	-0.607600	

```

4 -1.615544 -1.540707 -1.423087      -1.272086      -1.215968      -1.287337

    Shell weight      Age
0     -0.638217  1.577830
1     -1.212987 -0.919022
2     -0.207139 -0.294809
3     -0.602294  0.017298
4     -1.320757 -0.919022

```

Split the data into training and testing

```
X_train, X_test, Y_train, Y_test = train_test_split(indep_Scaled, dep,
test_size=0.5, random_state=0)
```

```
X_train.shape,X_test.shape
```

```
((2088, 8), (2089, 8))
```

```
Y_train.shape,Y_test.shape
```

```
((2088,), (2089,))
```

```
X_train.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	\
4149	-2.031938	-1.943819	-1.662195	-1.437280	-1.335371	
315	-0.616198	-0.532927	-0.705762	-0.755087	-0.745112	
1604	0.133312	0.474853	0.011563	0.292166	0.268692	
926	-0.782755	-0.835261	-1.183978	-0.823409	-0.641479	
3077	1.340855	1.029133	0.609334	1.460766	1.471739	

	Viscera weight	Shell weight	Age
4149	-1.451569	-1.500373	-1.231128
315	-1.036427	-0.710063	-0.606915
1604	0.678883	0.216754	0.017298
926	-0.853947	-1.033372	-0.919022
3077	1.796573	1.197456	0.017298

```
X_test.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	\
668	0.216591	0.172519	0.370226	0.181016	-0.368878	
1580	-0.199803	-0.079426	-0.466653	-0.433875	-0.443224	
3784	0.799543	0.726798	0.370226	0.870348	0.755318	
463	-2.531611	-2.447709	-2.020857	-1.579022	-1.522362	
2615	1.007740	0.928354	0.848442	1.390405	1.415417	

	Viscera weight	Shell weight	Age
668	0.569396	0.690940	0.953617
1580	-0.343004	-0.325685	-0.606915
3784	1.764639	0.565209	0.329404

```

463      -1.538247    -1.572219 -1.543234
2615      1.778325     0.996287  0.641511

Y_train.head()

4149    1
315     1
1604    2
926     1
3077    2
Name: Sex, dtype: int32

Y_test.head()

668     2
1580    1
3784    2
463     1
2615    2
Name: Sex, dtype: int32

```

Train the model

```

model = RandomForestClassifier(n_estimators=10,criterion='entropy')

model.fit(X_Train,Y_Train)

RandomForestClassifier(criterion='entropy', n_estimators=10)

y_predict = model.predict(X_test)

```

Error matrix

```

accuracy_score(Y_test,y_predict)

0.8085208233604595

pd.crosstab(Y_test,y_predict)

   col_0    0    1    2
Sex
0      521   23  108
1       39  613   41
2      121   68  555

print(classification_report(Y_test,y_predict))

          precision    recall  f1-score   support

           0       0.77      0.80      0.78      652
           1       0.87      0.88      0.88      693
           2       0.79      0.75      0.77      744

```

accuracy			0.81	2089
macro avg	0.81	0.81	0.81	2089
weighted avg	0.81	0.81	0.81	2089